



# Django Deployment Workshop

Jacob Kaplan-Moss  
Frank Wiles

**OSCON 2010**

<http://revsys.com/oscon2010>

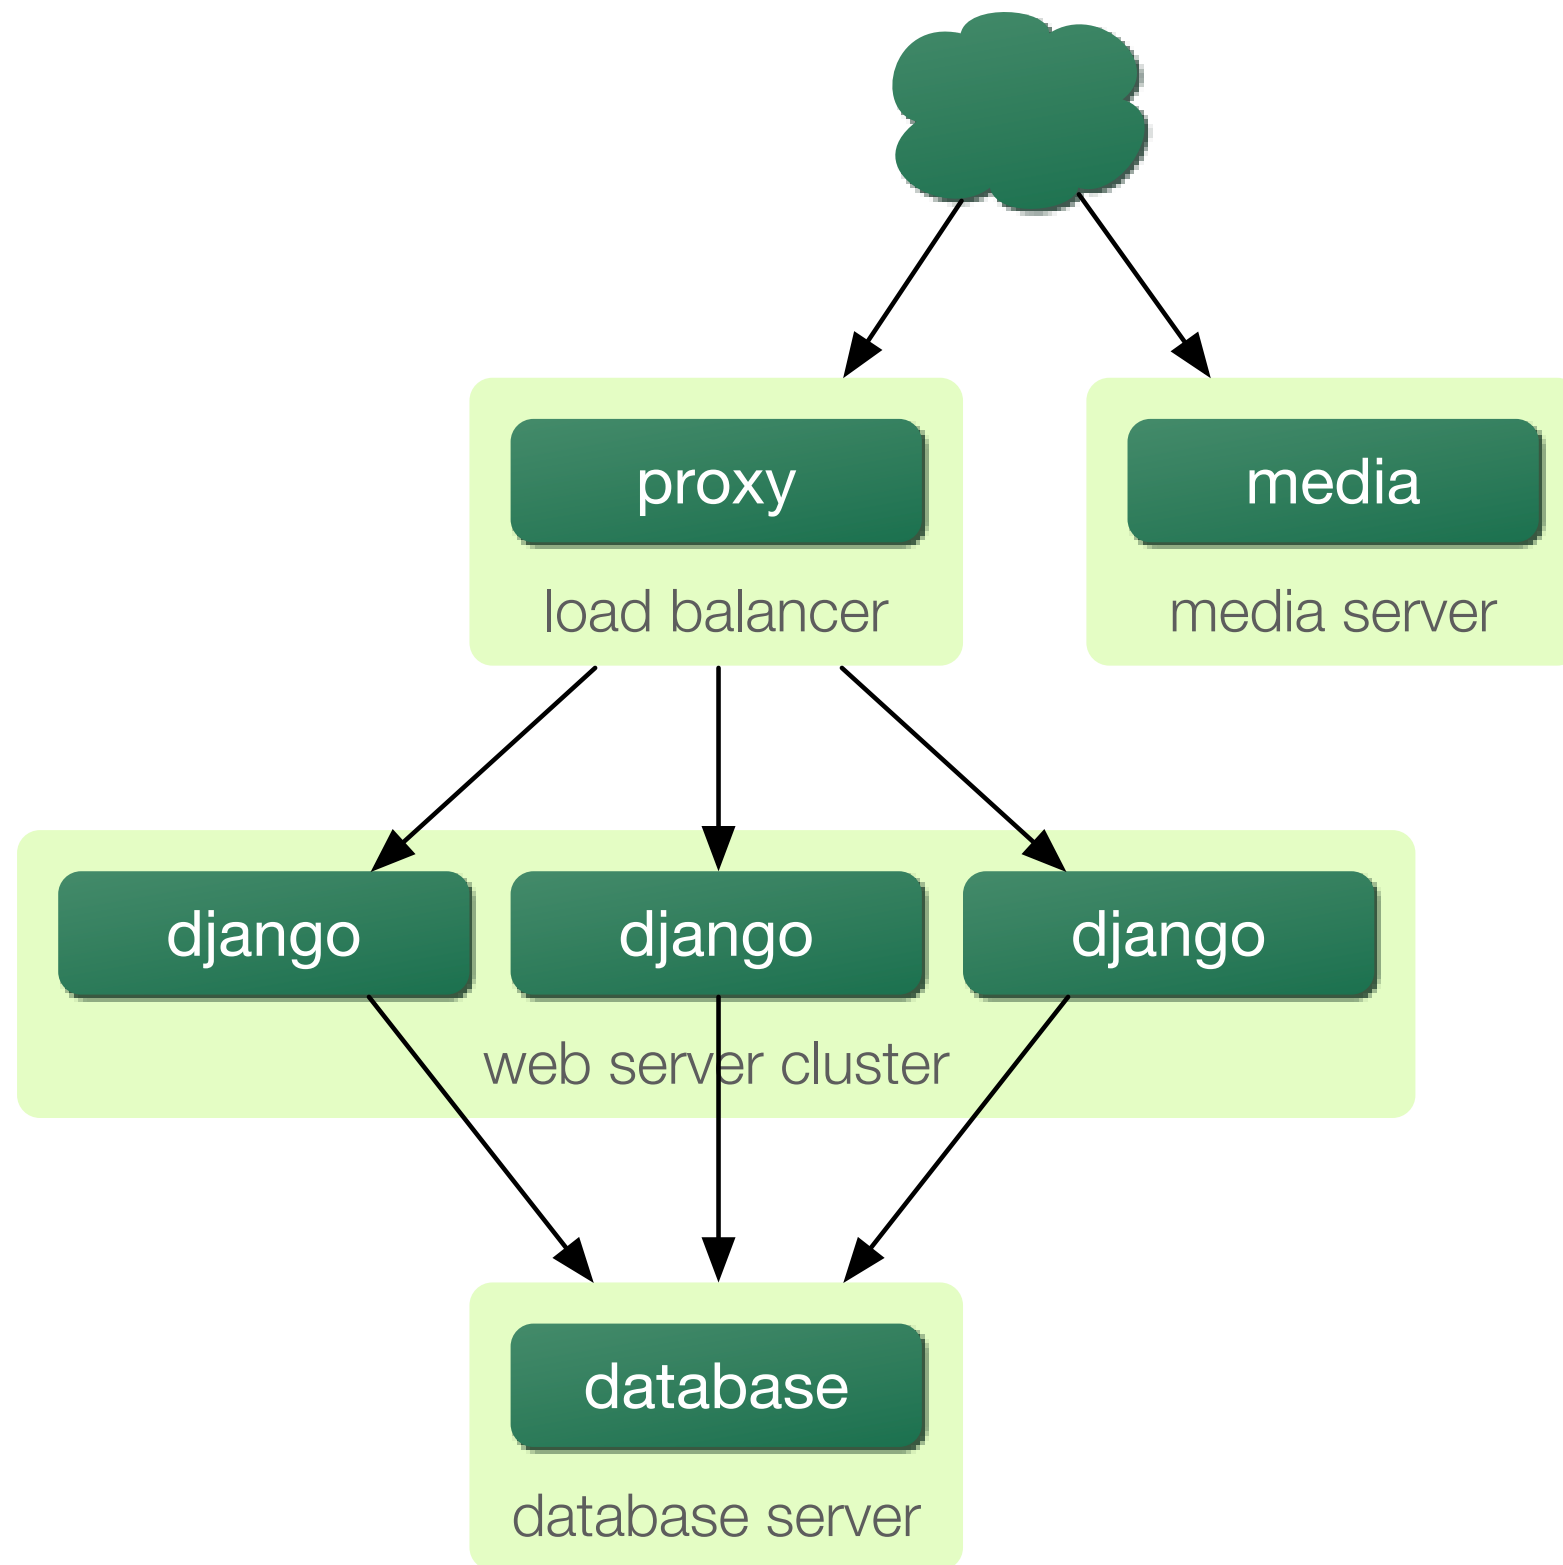
# About us

So you've written a  
Django site...

... now what?

- API Metering
- Backups & Snapshots
- Counters
- Cloud/Cluster Management Tools
  - Instrumentation/Monitoring
  - Failover
  - Node addition/removal and hashing
  - Autoscaling for cloud resources
- CSRF/XSS Protection
- Data Retention/Archival
- Deployment Tools
  - Multiple Devs, Staging, Prod
  - Data model upgrades
  - Rolling deployments
  - Multiple versions (selective beta)
  - Bucket Testing
  - Rollbacks
  - CDN Management
- Distributed File Storage
- Distributed Log storage, analysis
- Graphing
- HTTP Caching
- Input/Output Filtering
- Memory Caching
- Non-relational Key Stores
- Rate Limiting
- Relational Storage
- Queues
- Rate Limiting
- Real-time messaging (XMPP)
- Search
  - Ranging
  - Geo
- Sharding
- Smart Caching
  - Dirty-table management

# What we're building



# Deployment steps

1. Bootstrap an app instance.
2. Configure the database.
3. Set up application servers.
4. Automate deployment.
5. Scale up to multiple web servers.
6. Start caching.
7. Tune, tune, tune...



# Deployment steps

- 1. Bootstrap an app instance.**
2. Configure the database.
3. Set up application servers.
4. Automate deployment.
5. Scale up to multiple web servers.
6. Start caching.
7. Tune, tune, tune...

# Demo

# Deployment steps

1. Bootstrap an app instance.
- 2. Configure the database.**
3. Set up application servers.
4. Automate deployment.
5. Scale up to multiple web servers.
6. Start caching.
7. Tune, tune, tune...

# Database servers

- SQLite
- MySQL
- PostgreSQL
- Oracle
- ...

# Why PostgreSQL?

# Demo

# Deployment steps

1. Bootstrap an app instance.
2. Configure the database.
- 3. Set up application servers.**
4. Automate deployment.
5. Scale up to multiple web servers.
6. Start caching.
7. Tune, tune, tune...

# App servers should...

- Have proven reliability.
- Be highly stable.
- Have predictable resource consumption.
- Speak WSGI.



# Application servers

- Old school: Apache + mod\_python
- New school: Apache + mod\_wsgi
- Cutting edge: uWSGI, Gunicorn
- Avoid: FastCGI, SCGI, AJP, ...

Recommendation:  
Apache + mod\_wsgi

# Demo

# Deployment steps

1. Bootstrap an app instance.
2. Configure the database.
3. Set up application servers.
- 4. Automate deployment.**
5. Scale up to multiple web servers.
6. Start caching.
7. Tune, tune, tune...

# Deployment should...

- Automatically manage dependencies.
- Be isolated.
- Be automated.
- Be repeatable.
- Be identical in staging and in production.
- Work the same for everyone.

<b>Dependency management</b>	<b>Isolation</b>	<b>Automation</b>
apt/yum/...	virtualenv	Capistrano
easy_install	zc.buildout	Fabric
pip		Puppet/Chef
zc.buildout		

# Dependency management

- The Python ecosystem rocks!
- Python package management doesn't.
- Installing packages — and dependancies — correctly is a lot harder than it should be; most defaults are wrong.
- Here be dragons.

# Vendor packages

- APT, Yum, ...
- The good: familiar tools; stability; handles dependancies not on PyPI.
- The bad: small selection; not (very) portable; hard to supply user packages.
- The ugly: **installs packages system-wide.**



# easy\_install

- The good: multi-version packages.
- The bad: requires 'net connection; can't uninstall; can't handle non-PyPI packages; multi-version packages barely work.
- The ugly: stale; unsupported; defaults almost totally wrong; **installs system-wide.**

# pip

<http://pip.openplans.org/>

- “Pip Installs Packages”
- The good: Just Works™; handles non-PyPI packages (including direct from SCM); repeatable dependancies; integrates with virtualenv for isolation.
- The bad: still young; not yet bundled.
- The ugly: haven’t found it yet.

# zc.buildout

<http://buildout.org/>

- The good: incredibly flexible; handles any sort of dependency; repeatable builds; reusable “recipes;” good ecosystem; handles isolation, too.
- The bad: often cryptic, INI-style configuration file; confusing duplication of recipes; sometimes *too* flexible.
- The ugly: chronically undocumented.

# Package isolation

- Why?
  - Site A requires Foo v1.0; site B requires Foo v2.0.
  - You need to develop against multiple versions of dependancies.

# Package isolation tools

- Virtual machines (Xen, VMWare, EC2, ...)
- Multiple Python installations.
- “Virtual” Python installations.
  - **virtualenv**  
<http://pypi.python.org/pypi/virtualenv>
  - **zc.buildout**  
<http://buildout.org/>

# Automation

# Why automate?

- “I can’t push this fix to the servers until Alex gets back from lunch.”
- “Sorry, I can’t fix that. I’m new here.”
- “Oops, I just made the wrong version of our site live.”
- “It’s broken! What’d you do!?”

# Automation basics

- SSH is right out.
- Don't futz with the server. Write a recipe.
- Deploys should be idempotent.



# Capistrano

<http://capify.org/>

- The good: lots of features; good documentation; active community.
- The bad: stale development; very “opinionated” and Rails-oriented.

# Fabric

<http://fabfile.org/>

- The good: very simple; flexible; actively developed; Python.
- The bad: few high-level commands; not yet “1.0”.

# Configuration management

- CFEngine, Puppet, Chef, ...
- Will handle a *lot* more than code deployment!
- Almost a necessity past  $N = 20$  or so.

# Recommendations

● Pip, Virtualenv, and Fabric

■ Buildout and Fabric.

◆ Buildout, Fabric, Puppet.

◆◆ Utility computing, Fabric, Chef.

# Demo

# Deployment steps

1. Bootstrap an app instance.
2. Configure the database.
3. Set up application servers.
4. Automate deployment.
- 5. Scale up to multiple web servers.**
6. Start caching.
7. Tune, tune, tune...

# Deployment steps

## **5. Scale up to multiple web servers.**

A.  $N > 1$ .

B. Load balancers.

C. Database connection middleware

# Deployment steps

## **5. Scale up to multiple web servers.**

**A.  $N > 1$ .**

B. Load balancers.

C. Database connection middleware



# Why multiple servers?

- Eliminate resource contention.
- Easier to optimize for scarcity.
- $1 \rightarrow 2$  is much harder than  $2 \rightarrow N$

“Shared nothing”

```
BALANCE = None
```

```
def balance_sheet(request):  
    global BALANCE  
    if not BALANCE:  
        bank = Bank.objects.get(...)  
        BALANCE = bank.total_balance()  
    ...
```

# FAIL

Global variables are  
right out

```
from django.cache import cache

def balance_sheet(request):
    balance = cache.get('bank_balance')
    if not balance:
        bank = Bank.objects.get(...)
        balance = bank.total_balance()
        cache.set('bank_balance', balance)
```

...

# WIN

```
def generate_report(request):  
    report = get_the_report()  
    open('/tmp/report.txt', 'w').write(report)  
    return redirect(view_report)  
  
def view_report(request):  
    report = open('/tmp/report.txt').read()  
    return HttpResponse(report)
```

# FAIL

Filesystem?  
What filesystem?

# Dealing with media



# Demo

# Deployment steps

## **5. Scale up to multiple web servers.**

A.  $N > 1$ .

**B. Load balancers.**

C. Database connection middleware

# Why load balancers?

# Load balancer traits

- Low memory overhead.
- High concurrency.
- Hot failover.
- Other nifty features...

# Load balancers

- Apache + mod\_proxy
- perlbal
- nginx
- Varnish / Squid

# Recommendation: Nginx

# Demo

# Deployment steps

## **5. Scale up to multiple web servers.**

A.  $N > 1$ .

B. Load balancers.

**C. Database connection middleware**



# Connection middleware?

```
DATABASE_HOST = '10.0.0.100'
```

# FAIL

# Connection middleware

- Proxy between web and database layers
- Most implement hot fallover and connection pooling
  - Some also provide replication, load balancing, parallel queries, connection limiting, ...
- `DATABASE_HOST = '127.0.0.1'`

# Connection middleware

- PostgreSQL: pgpool (I, II), pgbouncer.
- MySQL: MySQL Proxy.
- Database-agnostic: sqlrelay.
- Oracle: ?

# pgpool vs. pgbouncer

# Demo

# Deployment steps

1. Bootstrap an app instance.
2. Configure the database.
3. Set up application servers.
4. Automate deployment.
5. Scale up to multiple web servers.
- 6. Start caching.**
7. Tune, tune, tune...

“

*Performance is mostly  
an exercise in caching  
and cache invalidation.*”

XXX or some such quote - three hard  
problems, maybe?



# Caching layers

- Low-level (Python API)
- Template fragment caching
- Per-view cache
- Whole site cache
- Upstream caches
- CDNs

# Cache backends

# Production

- Use memcached.
- Really.

# Other cache backends

- filesystem
- database
- local memory

# Demo



# Deployment steps

1. Bootstrap an app instance.
2. Configure the database.
3. Set up application servers.
4. Automate deployment.
5. Scale up to multiple web servers.
6. Start caching.
- 7. Tune, tune, tune...**

# What's next?

- Database redundancy / replication
- Monitoring (availability; capacity planning)
- Searching, queuing, and locking (oh my!)
- NoSQL.
- Scaling up and out.



# Thank you!

Us: Jacob Kaplan-Moss

Frank Wiles

This talk: <http://revsys.com/oscon2010>

Web: <http://revsys.com/>

Twitter: @revsys