



# Django Deployment Workshop

Jacob Kaplan-Moss

<http://bit.ly/ddw-slides>

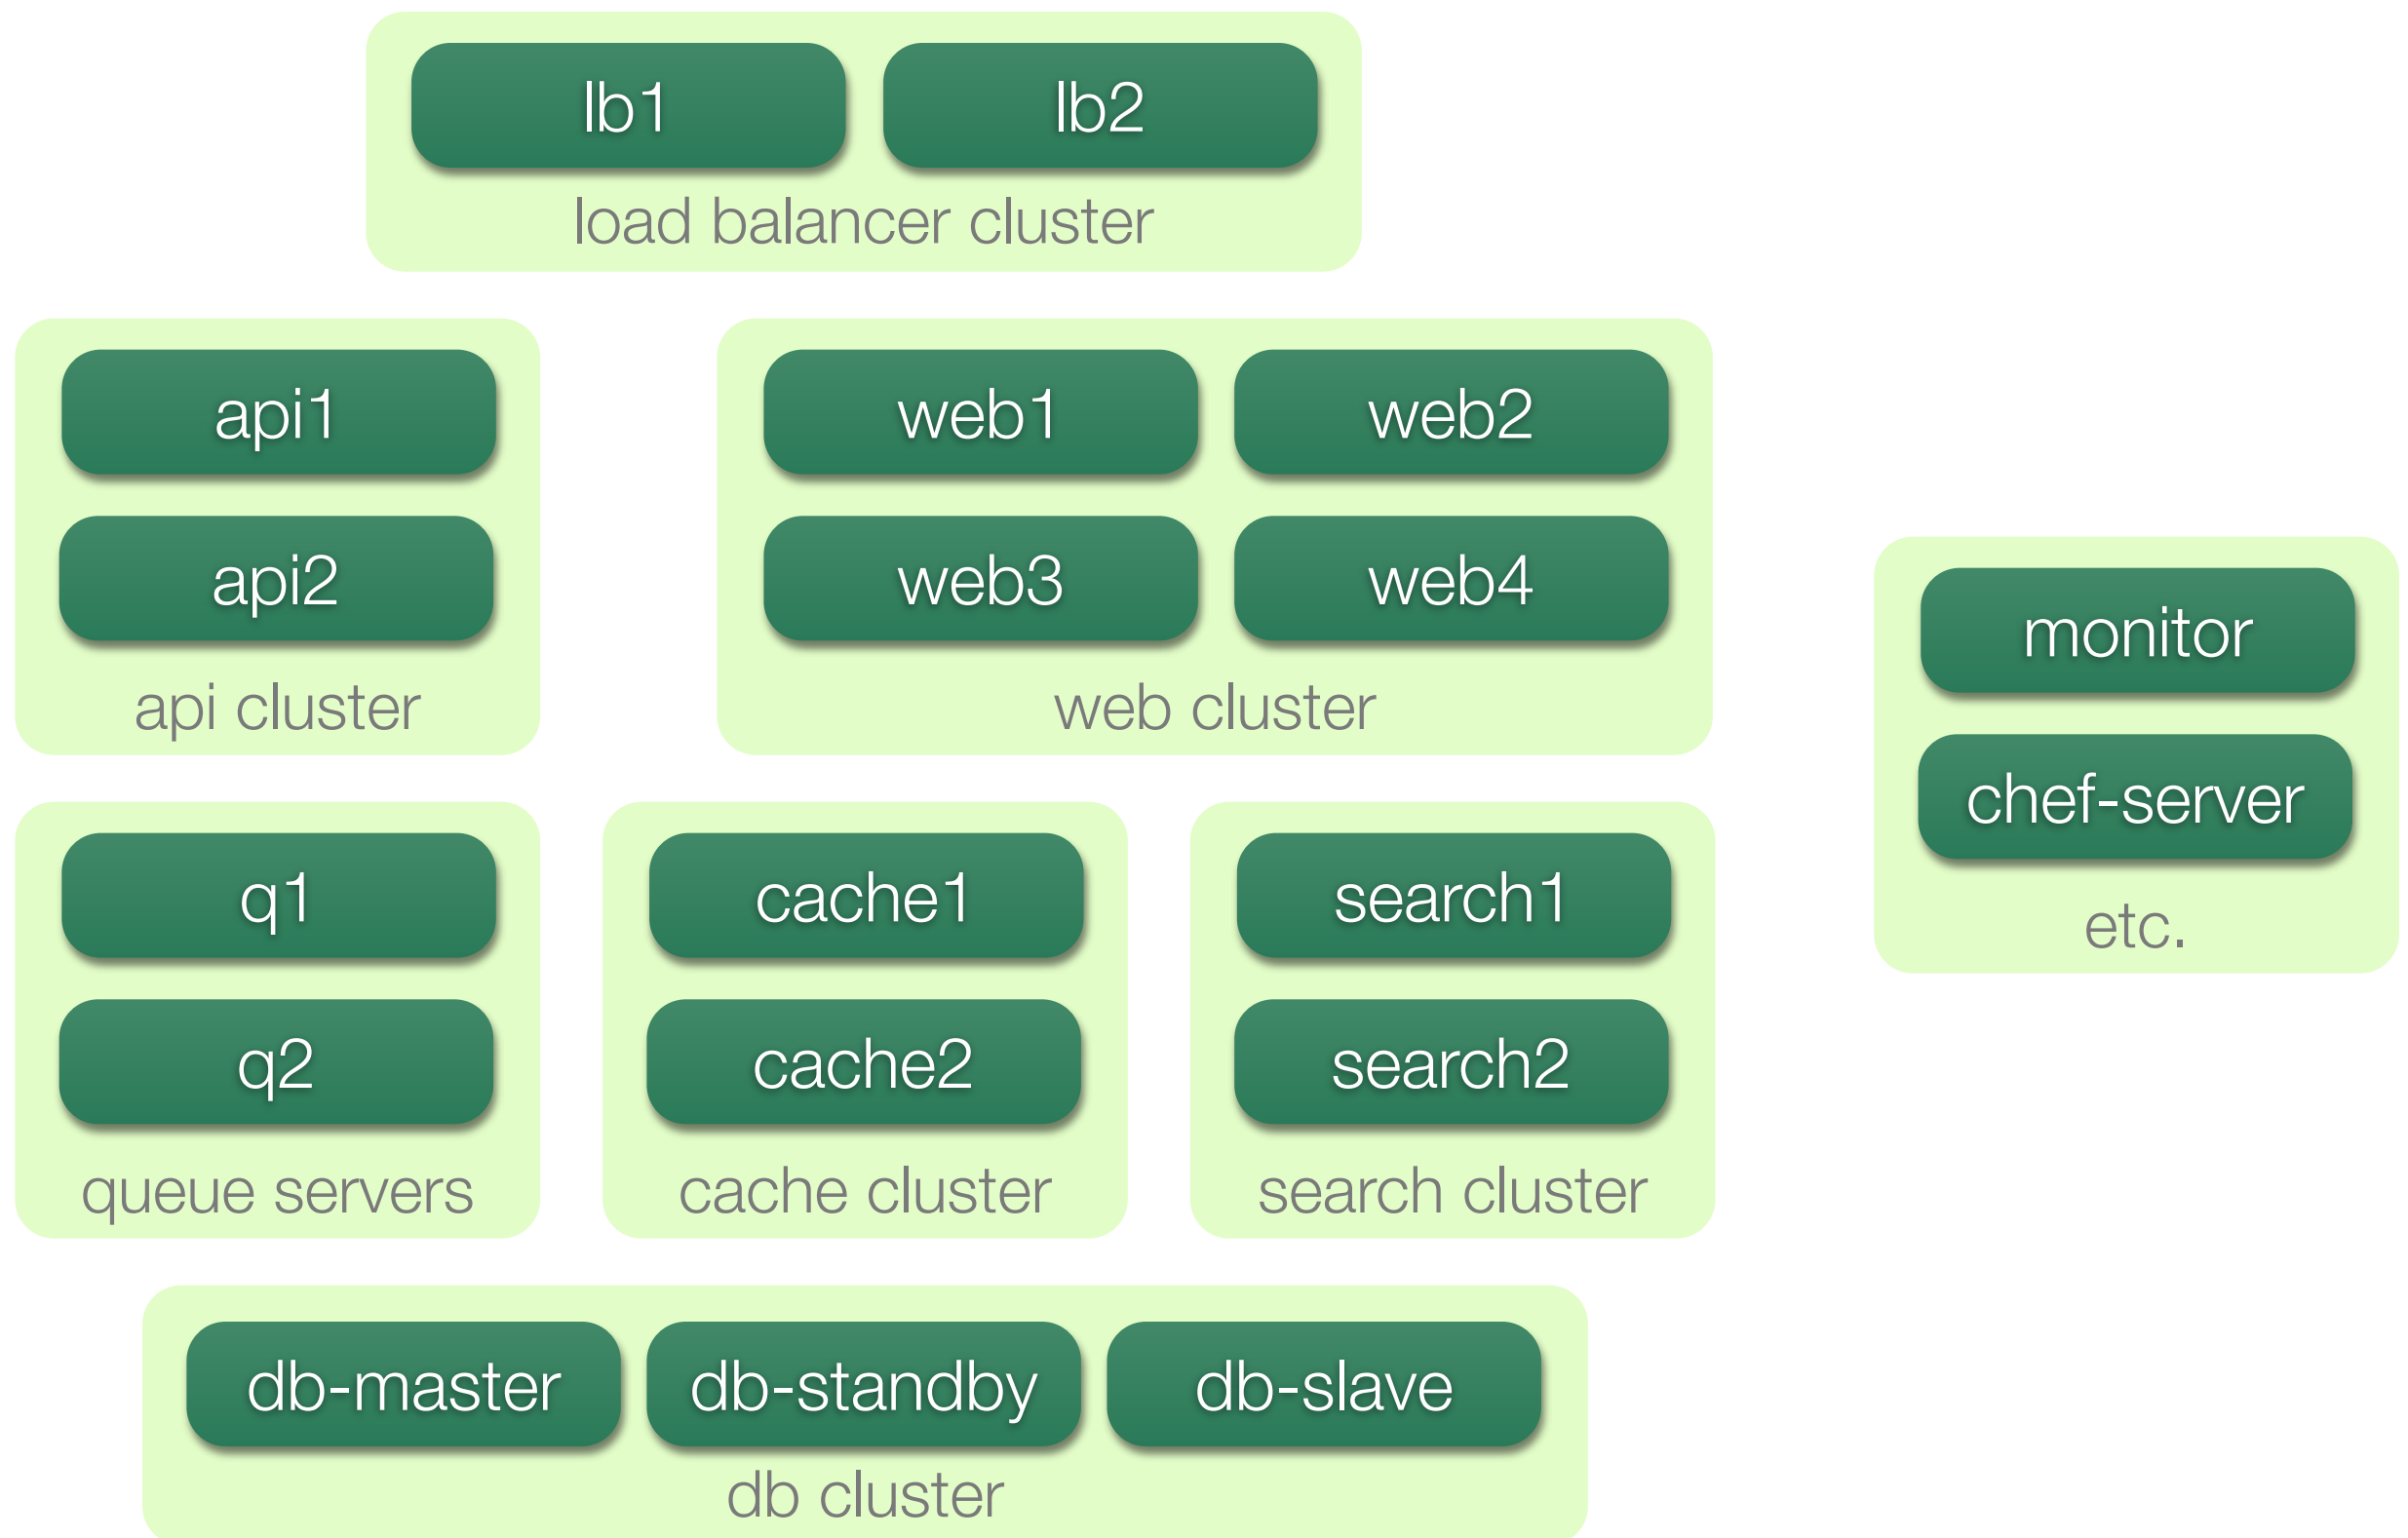
<http://bit.ly/ddw-code>

So you've written a  
Django site...

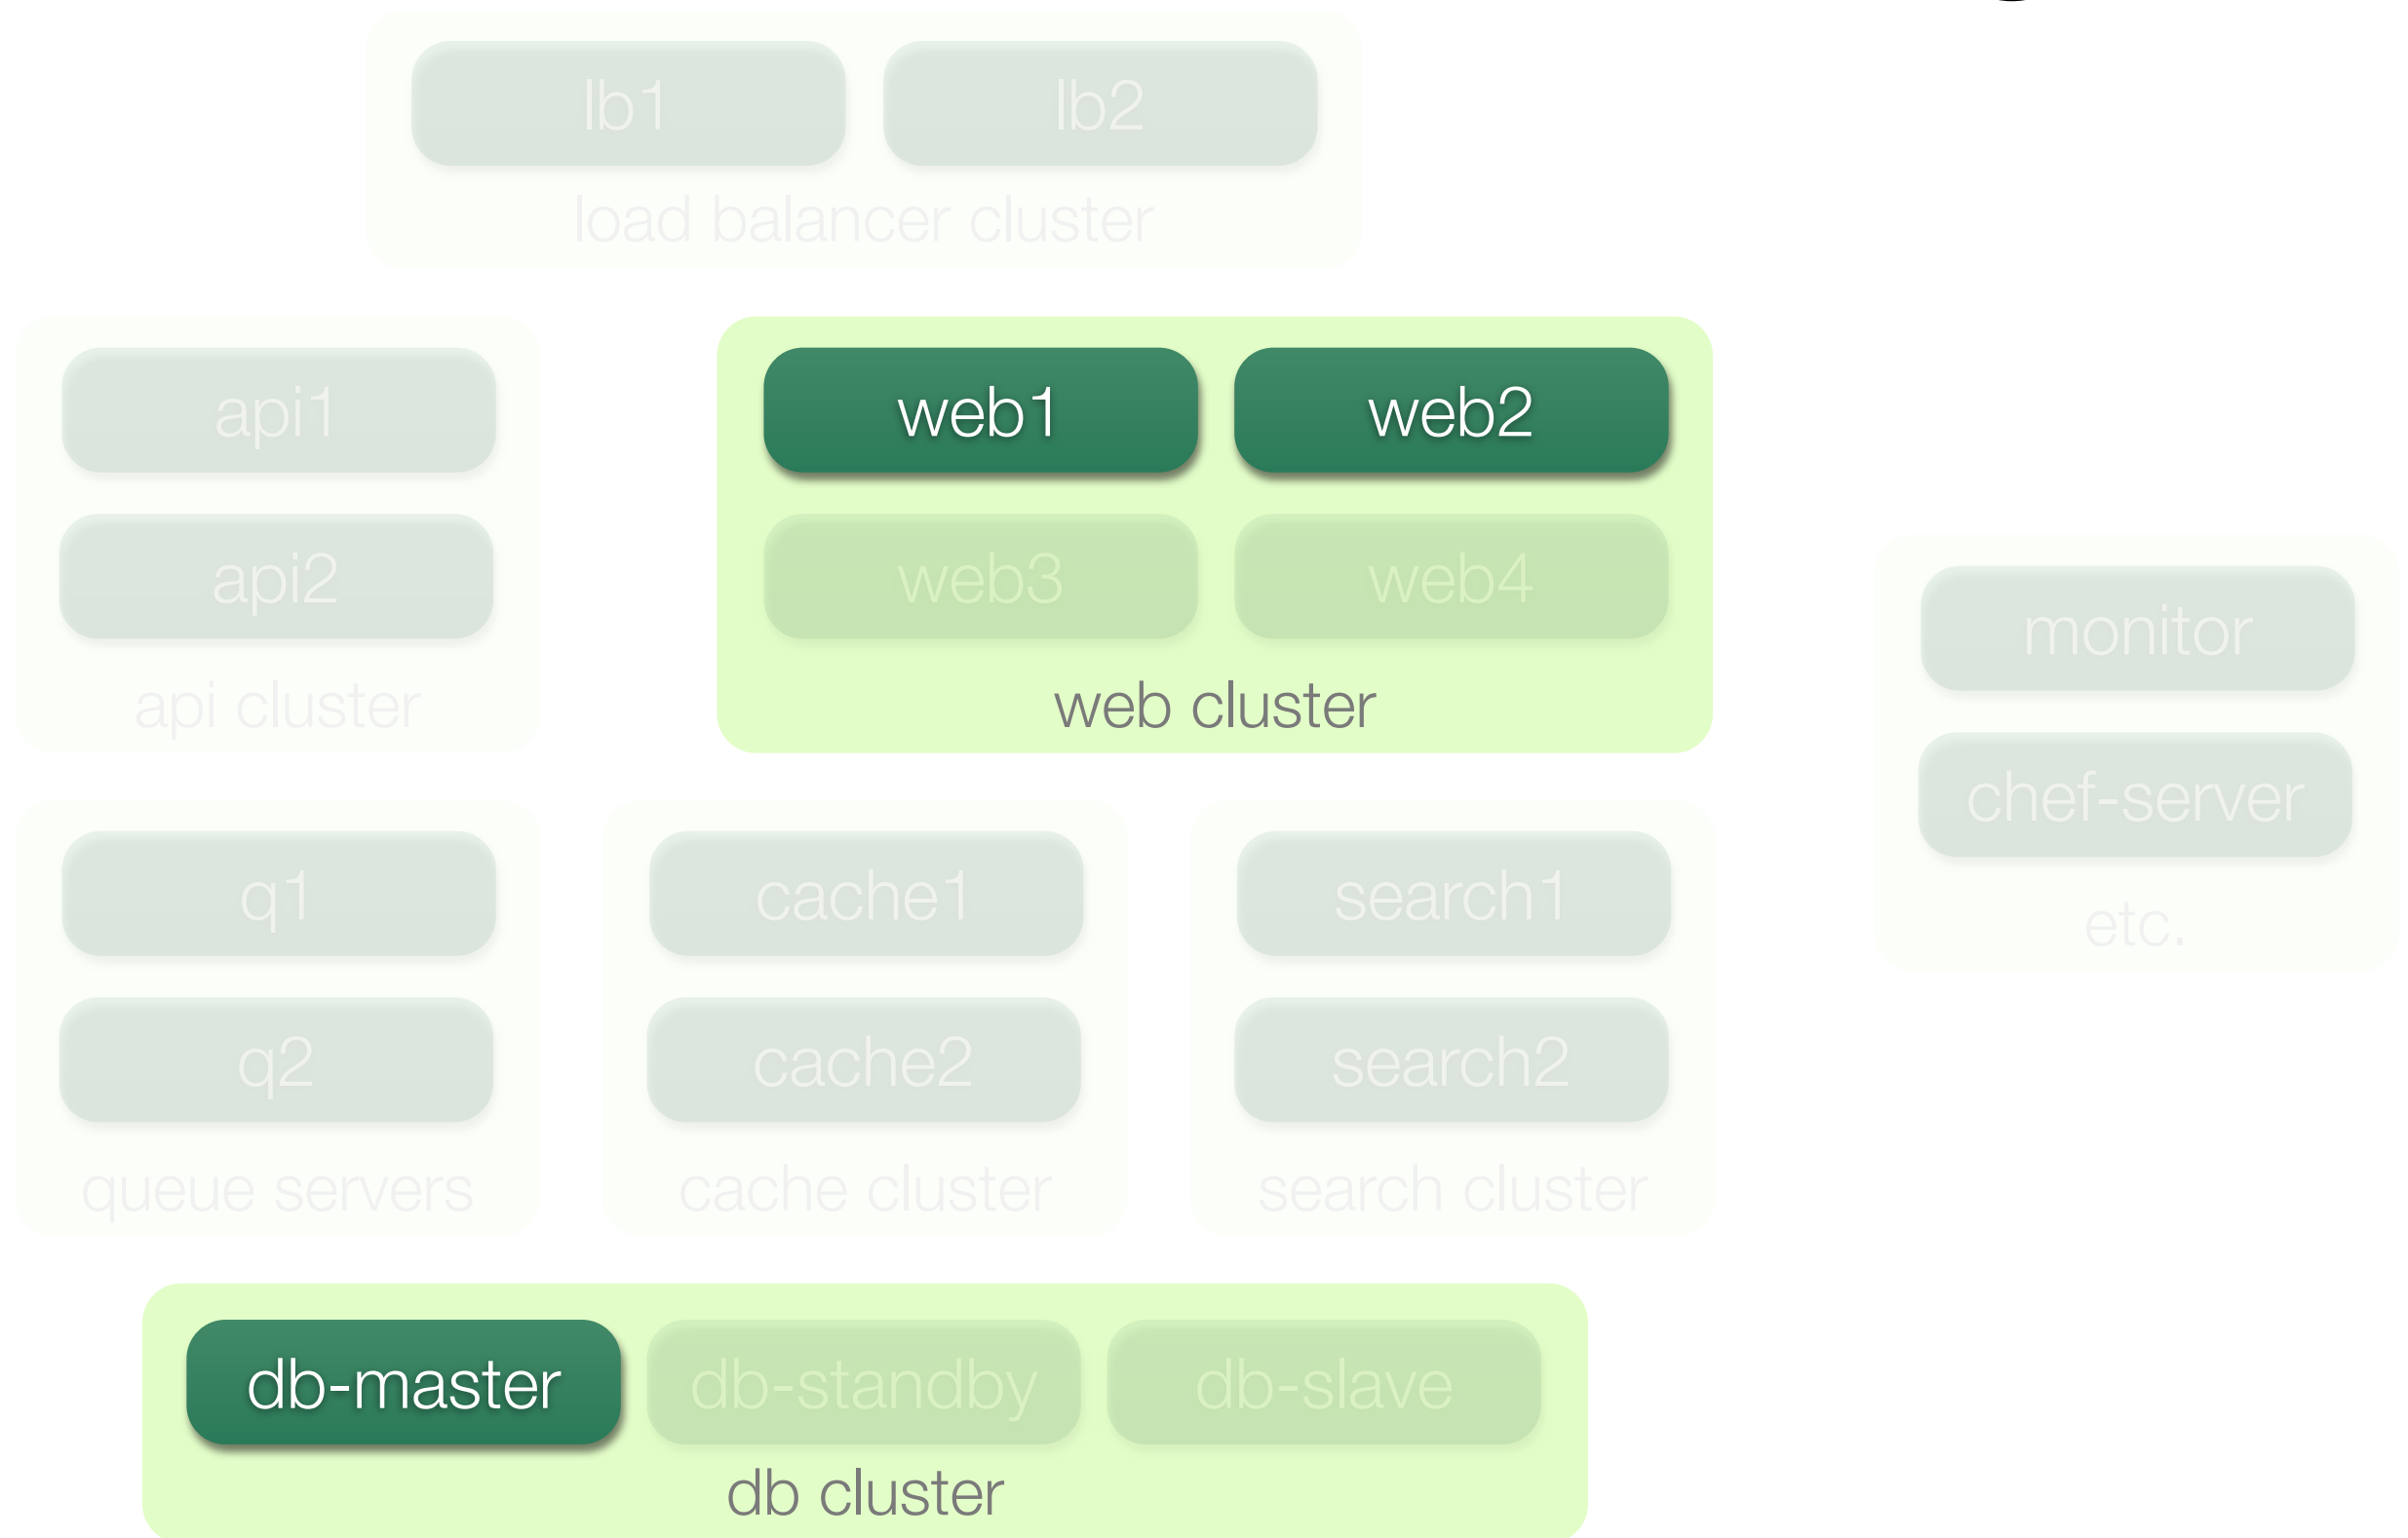
... now what?

- API Metering
- Backups & Snapshots
- Counters
- Cloud/Cluster Management Tools
  - Instrumentation/Monitoring
  - Failover
  - Node addition/removal and hashing
  - Autoscaling for cloud resources
- CSRF/XSS Protection
- Data Retention/Archival
- Deployment Tools
  - Multiple Devs, Staging, Prod
  - Data model upgrades
  - Rolling deployments
  - Multiple versions (selective beta)
  - Bucket Testing
  - Rollbacks
  - CDN Management
- Distributed File Storage
- Distributed Log storage, analysis
- Graphing
- HTTP Caching
- Input/Output Filtering
- Memory Caching
- Non-relational Key Stores
- Rate Limiting
- Relational Storage
- Queues
- Rate Limiting
- Real-time messaging (XMPP)
- Search
  - Ranging
  - Geo
- Sharding
- Smart Caching
  - Dirty-table management

# A real-world stack



# What we're building



# Deployment steps

1. Bootstrap an app instance.
2. Configure the database.
3. Set up application servers.
4. Automate deployment.
5. Scale up to multiple web servers.
6. Tune, tune, tune...

# Deployment steps

- 1. Bootstrap an app instance.**
2. Configure the database.
3. Set up application servers.
4. Automate deployment.
5. Scale up to multiple web servers.
6. Tune, tune, tune...



# Dependency management

- The Python ecosystem rocks!
- Python package management doesn't.
- Installing packages — and dependancies — correctly is a lot harder than it should be; most defaults are wrong.
- Here be dragons.

# Vendor packages

- APT, Yum, ...
- The good: familiar tools; stability; handles dependancies not on PyPI.
- The bad: small selection; not (very) portable; hard to supply user packages.
- The ugly: **installs packages system-wide.**

# easy\_install

- The good: multi-version packages.
- The bad: requires 'net connection; can't uninstall; can't handle non-PyPI packages; multi-version packages barely work.
- The ugly: stale; unsupported; defaults almost totally wrong; **installs system-wide.**

# pip

<http://pip.openplans.org/>

- “Pip Installs Packages”
- The good: Just Works™; handles non-PyPI packages (including direct from SCM); repeatable dependancies; integrates with virtualenv for isolation.
- The bad: still young; not yet bundled.
- The ugly: haven’t found it yet.

# zc.buildout

<http://buildout.org/>

- The good: incredibly flexible; handles any sort of dependency; repeatable builds; reusable “recipes;” good ecosystem; handles isolation, too.
- The bad: often cryptic, INI-style configuration file; confusing duplication of recipes; sometimes *too* flexible.
- The ugly: chronically undocumented.

# Package isolation

- Why?
  - Site A requires Foo v1.0; site B requires Foo v2.0.
  - You need to develop against multiple versions of dependencies.

# Package isolation tools

- Virtual machines (Xen, VMWare, EC2, ...)
- Multiple Python installations.
- “Virtual” Python installations.
  - **virtualenv**  
<http://pypi.python.org/pypi/virtualenv>
  - **zc.buildout**  
<http://buildout.org/>

Suggestions:  
pip + virtualenv



# Demo

Bootstrapping the application.

# Deployment steps

1. Bootstrap an app instance.
- 2. Configure the database.**
3. Set up application servers.
4. Automate deployment.
5. Scale up to multiple web servers.
6. Tune, tune, tune...

# Database servers

- SQLite
- MySQL
- PostgreSQL
- Oracle
- ...

# Why PostgreSQL?

# Demo

Setting up PostgreSQL.

# Deployment steps

1. Bootstrap an app instance.
2. Configure the database.
- 3. Set up application servers.**
4. Automate deployment.
5. Scale up to multiple web servers.
6. Tune, tune, tune...

# App servers should...

- Have proven reliability.
- Be highly stable.
- Have predictable resource consumption.
- Speak WSGI.

# Application servers

- Outdated: Apache + mod\_python
- The new standard: Apache + mod\_wsgi
- Cutting edge: uWSGI, Gunicorn, CherryPy...
- Avoid: FastCGI, SCGI, AJP.



# Recommendations:

- **Most cases: Apache + mod\_wsgi.**


“A rule of thumb that has worked well for me is that if I'm excited to play around with something, it probably doesn't belong in production.” — Maciej Ceglowski

<http://pinboard.in/blog/63/>

- **Easier setup: Gunicorn.**

But beware: scaling considerations are still an unknown.

# Performance?



<http://nichol.as/benchmark-of-python-web-servers>

# Server performance doesn't matter unless your app is this simple.

```
def application(environ, start_response):  
    status = '200 OK'  
    output = 'Pong!'  
    response_headers = [('Content-type', 'text/plain'),  
                        ('Content-Length', str(len(output)))]  
    start_response(status, response_headers)  
    return [output]
```

# Demo

Playing with Apache, mod\_wsgi, and Gunicorn

# Deployment steps

1. Bootstrap an app instance.
2. Configure the database.
3. Set up application servers.
- 4. Automate deployment.**
5. Scale up to multiple web servers.
6. Tune, tune, tune...

# Deployment should...

- Automatically manage dependencies.
- Be isolated.
- Be automated.
- Be repeatable.
- Be identical in staging and in production.
- Work the same for everyone.

<b>Dependency management</b>	<b>Isolation</b>	<b>Automation</b>
apt/yum/...	virtualenv	Capistrano
easy_install	zc.buildout	Fabric
pip		
zc.buildout		

# Automation



# Why automate?

- “I can’t push this fix to the servers until Alex gets back from lunch.”
- “Sorry, I can’t fix that. I’m new here.”
- “Oops, I just made the wrong version of our site live.”
- “It’s broken! What’d you do!?”

# Automation basics

- SSH is right out.
- Don't futz with the server. Write a recipe.
- Deploys should be idempotent.

# Capistrano

<http://capify.org/>

- The good: lots of features; good documentation; active community.
- The bad: stale development; very “opinionated” and Rails-oriented.

# Fabric

<http://fabfile.org/>

- The good: very simple; flexible; actively developed; Python.
- The bad: few high-level commands;  
~~not yet “1.0”.~~

# Demo

A first fabfile.

# Deployment steps

1. Bootstrap an app instance.
2. Configure the database.
3. Set up application servers.
4. Automate deployment.
- 5. Scale up to multiple web servers.**
6. Tune, tune, tune...

# Deployment steps

## **5. Scale up to multiple web servers.**

A.  $N > 1$ .

B. Load balancers.

C. Database connection middleware

# Deployment steps

## **5. Scale up to multiple web servers.**

**A.  $N > 1$ .**

B. Load balancers.

C. Database connection middleware



# Why multiple servers?

- Eliminate resource contention.
- Easier to optimize for scarcity.
- $1 \rightarrow 2$  is much harder than  $2 \rightarrow N$

“Shared nothing”

```
BALANCE = None
```

```
def balance_sheet(request):  
    global BALANCE  
    if not BALANCE:  
        bank = Bank.objects.get(...)  
        BALANCE = bank.total_balance()  
    ...
```

# FAIL

Global variables are  
right out

```
from django.cache import cache

def balance_sheet(request):
    balance = cache.get('bank_balance')
    if not balance:
        bank = Bank.objects.get(...)
        balance = bank.total_balance()
        cache.set('bank_balance', balance)
```

...

# WIN

```
def generate_report(request):  
    report = get_the_report()  
    open('/tmp/report.txt', 'w').write(report)  
    return redirect(view_report)  
  
def view_report(request):  
    report = open('/tmp/report.txt').read()  
    return HttpResponse(report)
```

# FAIL

Filesystem?  
What filesystem?

# Dealing with media



# Demo

Rolling out a second web server.

# Deployment steps

## **5. Scale up to multiple web servers.**

A.  $N > 1$ .

**B. Load balancers.**

C. Database connection middleware

# Why load balancers?

# Load balancer traits

- Low memory overhead.
- High concurrency.
- Hot failover.
- Other nifty features...

# Load balancers

- Apache + mod\_proxy
- perlbal
- nginx
- Varnish / Squid

# Recommendation: Nginx

# Demo

# Deployment steps

## **5. Scale up to multiple web servers.**

A.  $N > 1$ .

B. Load balancers.

**C. Database connection middleware**



# Connection middleware?

```
DATABASE_HOST = '10.0.0.100'
```

# FAIL

# Connection middleware

- Proxy between web and database layers
- Most implement hot fallover and connection pooling
  - Some also provide replication, load balancing, parallel queries, connection limiting, ...
- `DATABASE_HOST = '127.0.0.1'`

# Connection middleware

- PostgreSQL: pgpool (I, II), pgbouncer.
- MySQL: MySQL Proxy.
- Database-agnostic: sqlrelay.
- Oracle: ?

# Deployment steps

1. Bootstrap an app instance.
2. Configure the database.
3. Set up application servers.
4. Automate deployment.
5. Scale up to multiple web servers.
- 6. Tune, tune, tune...**

# What's next?

- Caching.
- Configuration management.
- Database redundancy / replication.
- Monitoring (availability; capacity planning).
- Searching, queuing.
- Scaling up and out.

# Our toolkit:

- **Caching: memcached**

<http://memcached.org/>, <http://django.me/cache>

- **Config. management: Chef, Puppet.**

<http://www.opscode.com/chef/>, <http://www.puppetlabs.com/>

- **DB replication: PostgreSQL hot standby.**

[http://wiki.postgresql.org/wiki/Hot\\_Standby](http://wiki.postgresql.org/wiki/Hot_Standby)

- **Monitoring: Nagios, Munin.**

<http://www.nagios.org/>, <http://munin-monitoring.org/>

- **Searching: Haystack, Solr.**

<http://haystacksearch.org/>, <http://lucene.apache.org/solr/>

- **Queuing: Celery, Redis.**

<http://celeryproject.org/>, <http://redis.io/>

# Thank you!

Us: Jacob Kaplan-Moss

Frank Wiles

Web: <http://revsys.com/>

Twitter: @revsys



# Bonus:

# Configuration management

# The basics

- Automate your entire infrastructure.
- No more `ssh server`; `vi /etc/hosts`!
- Smart updates, self-healing, etc.
- Imperative rather than declarative.
- Almost a necessity past  $N \cong 20$  or so, but highly recommended even for  $N = 1$ .

# Options

- **CFEngine**

The original, still somewhat popular, but fairly crusty.

- **Bcfg2**

Python - yay! XML - boo!

- **Puppet**

The first fairly popular modern tool. Probably the most popular among LAMPish stacks.

- **Chef**

The new upstart, gaining lots of steam.

Our suggestion:  
Chef or Puppet

# Demo: Chef